



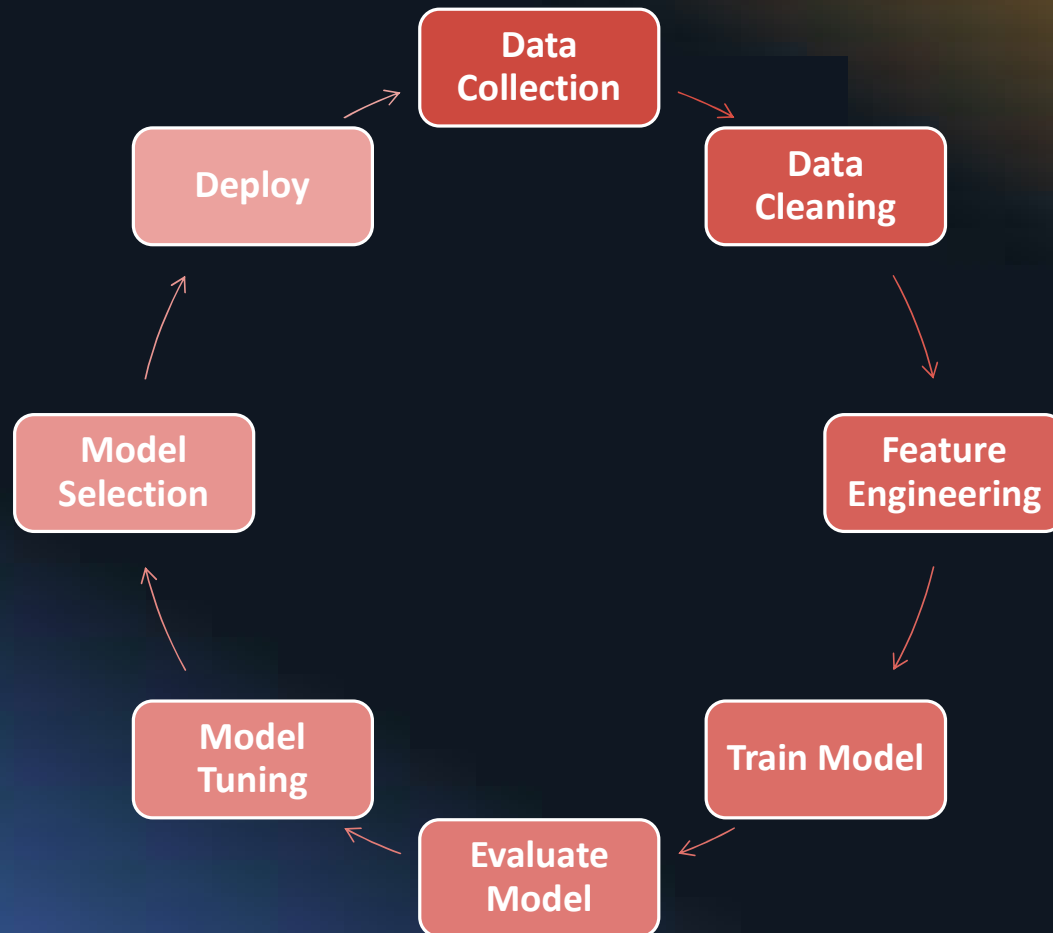
BUILD AN EXPLAINABLE AI MODEL WITH AWS SAGEMAKER STUDIO

SARBANI MAITI

VP DATA SCIENCE AT VEERSA TECHNOLOGY

EMAIL: SARBANIITB2020@GMAIL.COM

Machine Learning Lifecycle



MACHINE LEARNING LIFECYCLE IN PRODUCTION

Once the model is deployed, we need a production ready solution to serve the consumer applications.

Availability

- Availability of the ML system is critical for production system
- Onboarding the right kind of tool and technologies which are easy to adopt as per the change management system of the organization.
- Maintenance and reusability of the tools and models.
- Tracking, monitoring, alerting and feedback loop are other important aspects of the model in production


Scalability

- Model must support automatic scaling in production. Autoscaling dynamically adjusts the number of instances provisioned for a model in response to changes in your workload.
- When the workload increases, autoscaling brings more instances online.
- When the workload decreases, autoscaling removes unnecessary instances so that users don't pay for provisioned instances that are not in use.

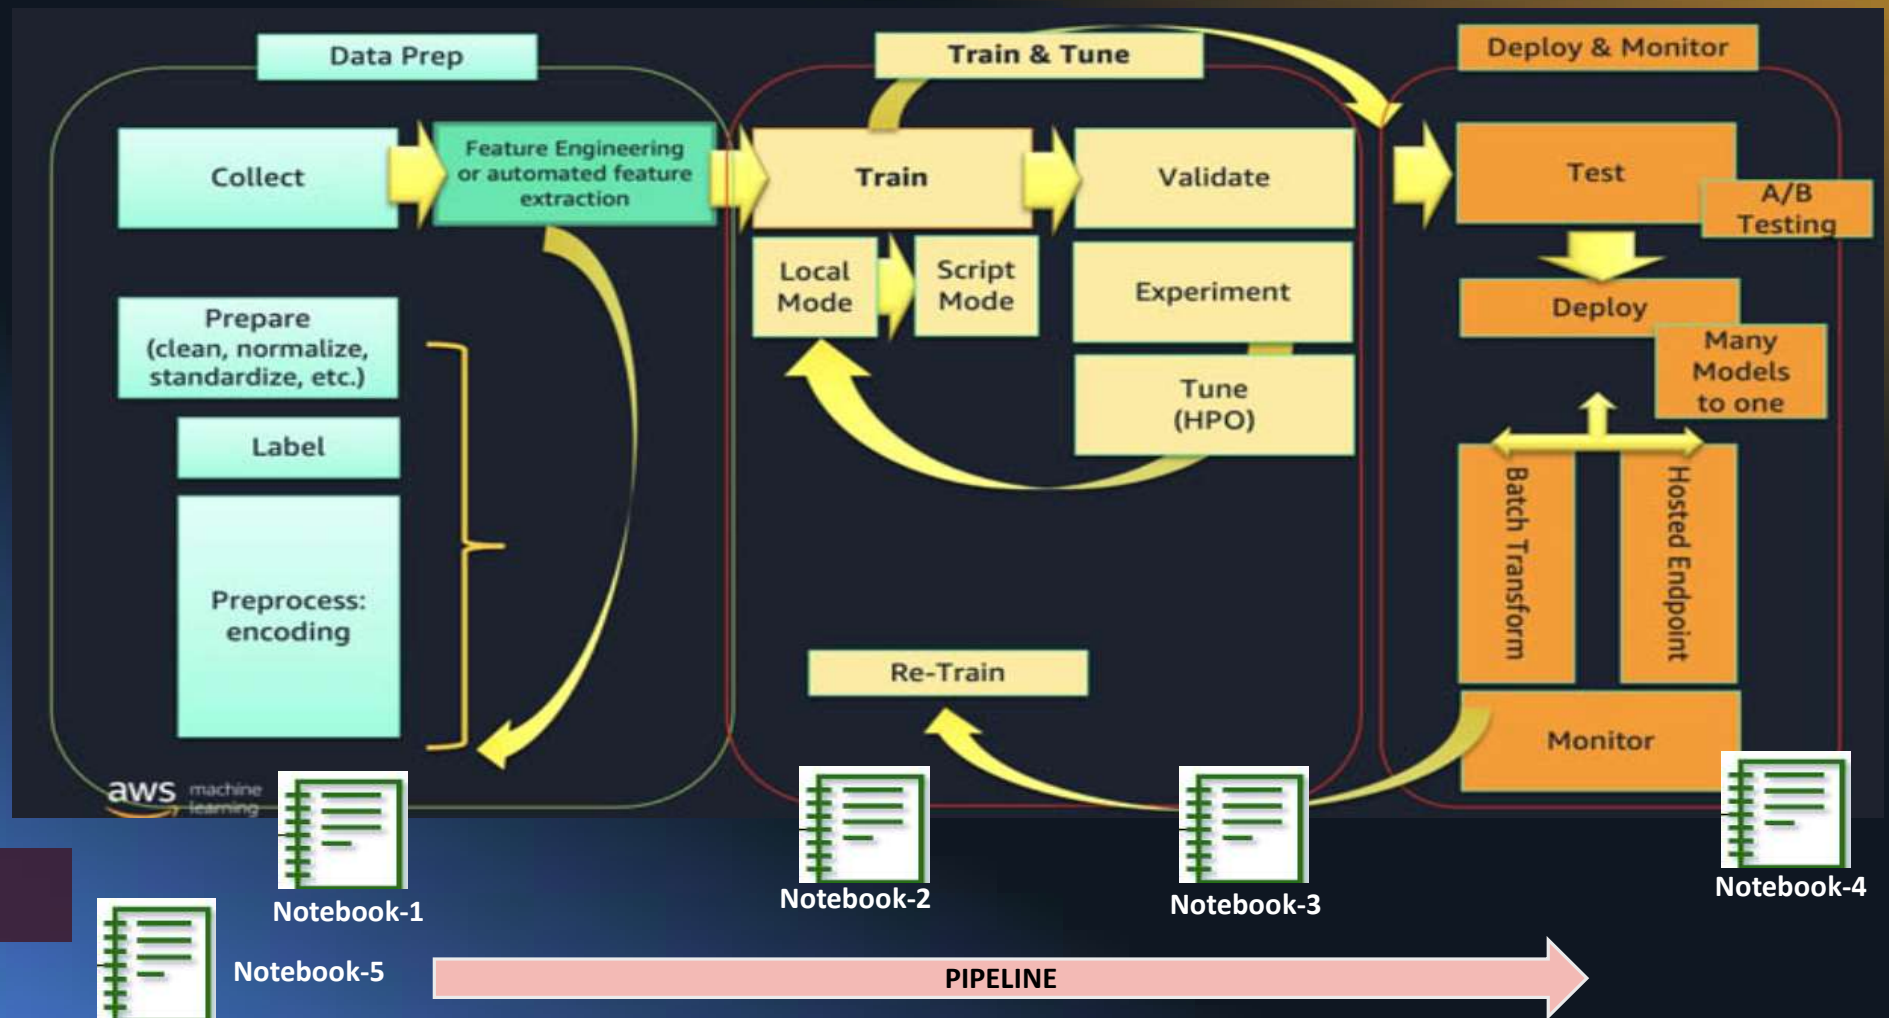
Security

- The model has to be compliant as per the regulatory requirement.
- Create environments with the least privileged access to sensitive data
- Protect & Encrypt sensitive data
- Audit and trace activity in your environment
- Reproduce results in your environment by tracking the lineage of ML artifacts throughout the lifecycle and using source and version control tools

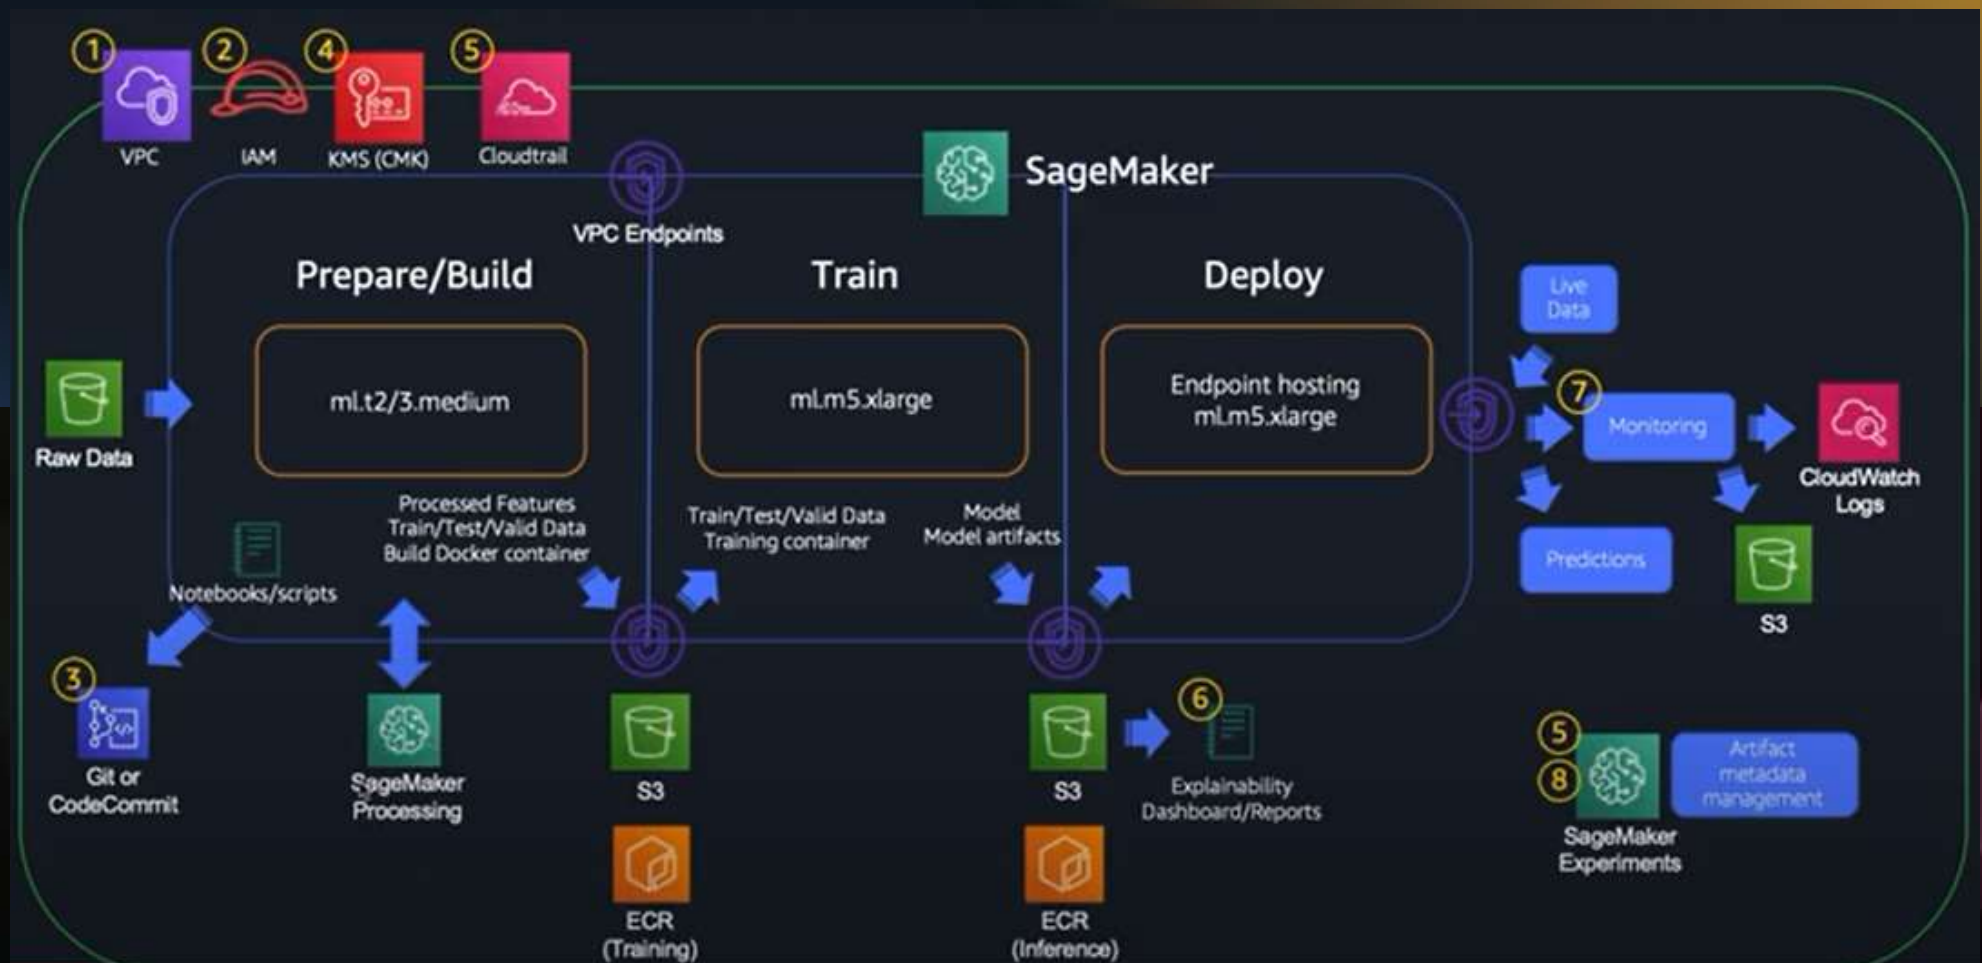
SECURE & COMPLAINT ML WORKFLOW

- 1 **COMPUTE & NETWORK ISOLATION** Deploy SageMaker in a VPC with no Internet access
- 2 **AUTHENTICATION & AUTHORIZATION** Provide single user access to Jupyter over IAM
- 3 **ARTIFACT MANAGEMENT** Enable private Git integration, lifecycle config, and versioning
- 4 **DATA ENCRYPTION** Encrypt data at motion and at rest across all ML workflow
- 5 **TRACEABILITY & AUDITABILITY** Trace model lineage, and audit all API calls and data events
- 6 **EXPLAINABILITY & INTERPRETABILITY** Explain predictions with feature importance and SHAP values
- 7 **REAL-TIME MODEL MONITORING** Monitor the performance of a productionized model
- 8 **REPRODUCIBILITY** Reproduce the model and results based on saved artifacts 

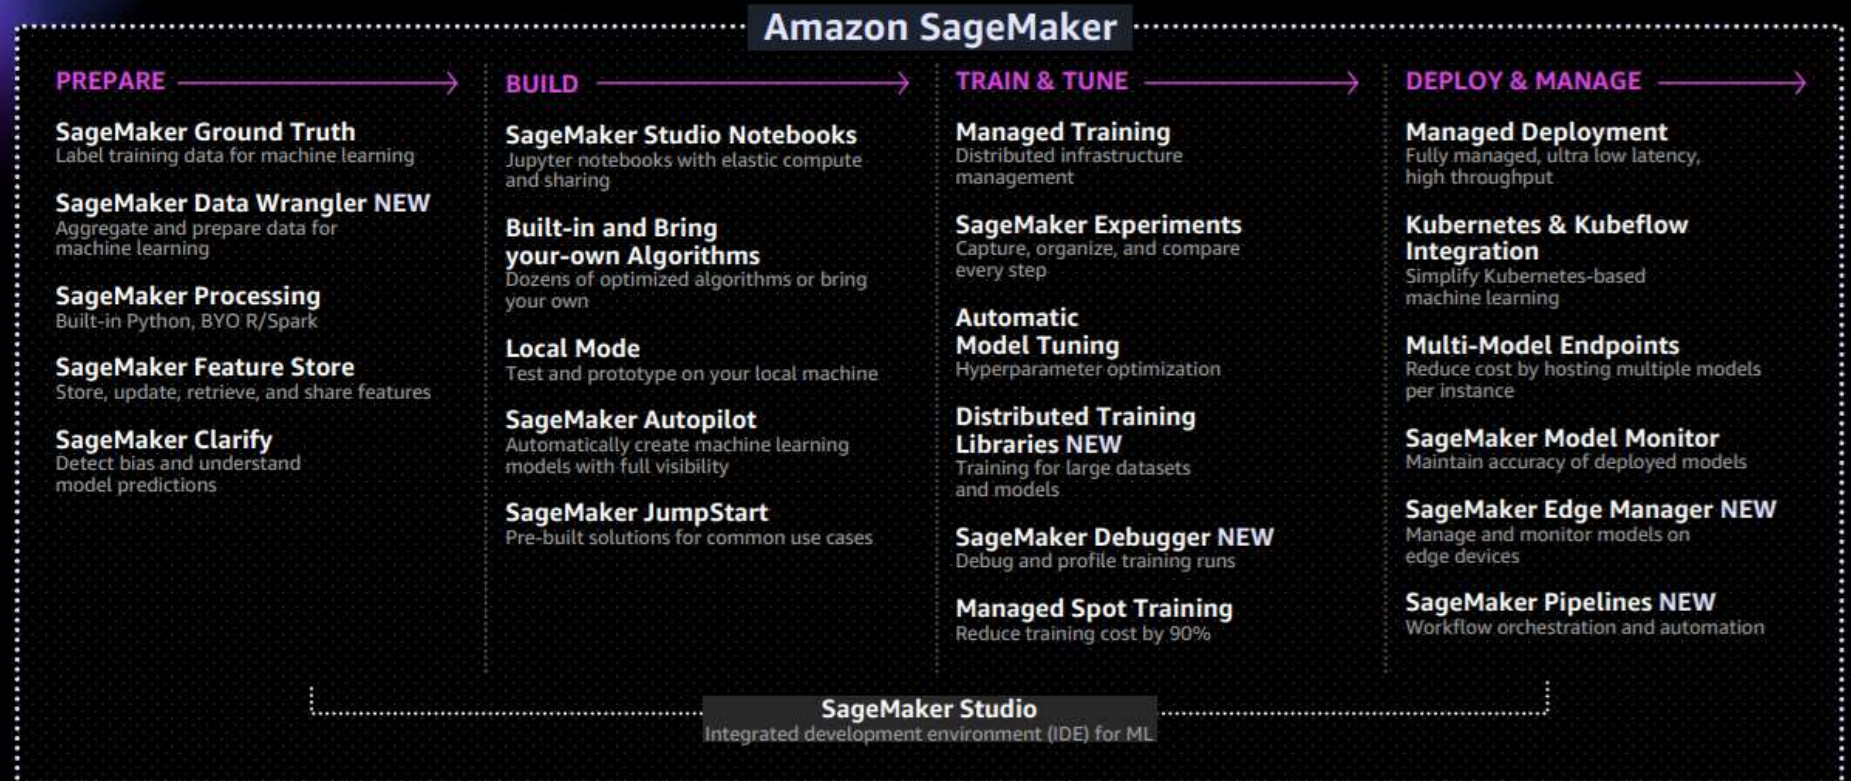
REFERENCE ML SYSTEM



SAGEMAKER SOLUTION



Amazon SageMaker overview



Build ML models : full to no customization

CUSTOM BUILD MODELS

- Your own machine learning code
- Opensource/ proprietary implementations or frameworks



- More time to build
- More costly
- High level of ML expertise
- **Most flexible**

Build ML models : full to no customization

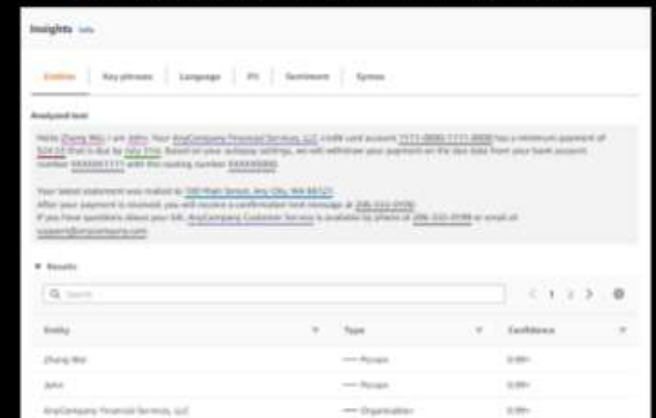
PRE-TRAINED MODELS: WHY TRAIN WHEN YOU CAN FINE TUNE?

Pre-trained model: a model that was trained on a large benchmark dataset to solve a problem similar to the one you want to solve.

- Use the deployed version via an API call (managed)
- Deploy directly (self-managed)
- Fine tune (Managed or self managed)

- Saves time
- Saves experimentation and cost
- ML expertise
- None to some flexibility

Amazon Comprehend real time text analysis



The screenshot shows the Amazon Comprehend console. At the top, there's a navigation bar with 'Insights' selected. Below it, there's a section for 'Analyzed text' showing a sample of text from a document. The main part of the interface is a table titled 'Results' which lists detected entities. The table has columns for 'Entity', 'Type', and 'Confidence'.

Entity	Type	Confidence
Zhang Wei	Person	0.99%
John	Person	0.99%
Amazon.com Services, LLC	Organization	0.99%

Amazon Rekognition label detection



Amazon SageMaker Studio Notebook

Perform data engineering,
analytics, and ML workflows
in one notebook



**Connect with Amazon EMR,
Amazon S3, and more**



**Interactively access, transform, and
analyze a wide range of data**



**Build, train, and deploy models using
your preferred framework**

Model options

AMAZON SAGEMAKER



Training code



No code required

- **XGBoost**
- **Matrix Factorization**
- **Regression**
- **Principal Component Analysis**
- **K-Means Clustering**
- **And More!**

Built-in Algorithms (17)
No ML coding required



Bring your Own Script
Amazon SageMaker build the container
Open source containers



Bring your Own Container
Full control, you build the container
R, C++, etc



Amazon SageMaker
Autopilot



Amazon SageMaker
Jumpstart

Fully Managed, Distributed, Auto-Scaled and Secure

Amazon SageMaker in-built algorithms

Amazon SageMaker
has built-in algorithms
or bring your own

Classification

Linear Learner | XGBoost | KNN

Computer vision

Image classification | Object detection |
Semantic segmentation

Topic modeling

LDA | NTM

Recommendation

Factorization machines

Forecasting

DeepAR

Working with text

BlazingText | Supervised | Unsupervised

Regression

Linear Learner | XGBoost | KNN

Clustering

KMeans

Sequence translation

Seq2Seq

Anomaly detection

Random cut forests | IP Insights

Feature reduction

PCA



Amazon SageMaker in-built algorithms

SAMPLE CODE

```
region=boto3.Session().region_name
container = sagemaker.image_uris.retrieve(region=region, framework='xgboost', version='latest')
print( "Using SageMaker XGBoost container: {}, ({}).format (container, region))

using SageMaker XGBoost container: 544295431143.dkr.ecr.ap-southeast-2.amazonaws.com/xgboost:latest, (ap-southeast-2)
```

```
sess = sagemaker.Session()

xgb = sagemaker.estimator.Estimator(container,
                                    role,
                                    instance_count=1,
                                    instance_type='ml.m4.xlarge',
                                    output_path='s3://{}/{}/output'.format(bucket, prefix),
                                    sagemaker_session=sess)

xgb.set_hyperparameters(max_depth=5,
                        eta=0.2,
                        gamma=4,
                        min_child_weight=6,
                        subsample=0.8,
                        silent=0,
                        objective='binary:logistic',
                        num_round=100)

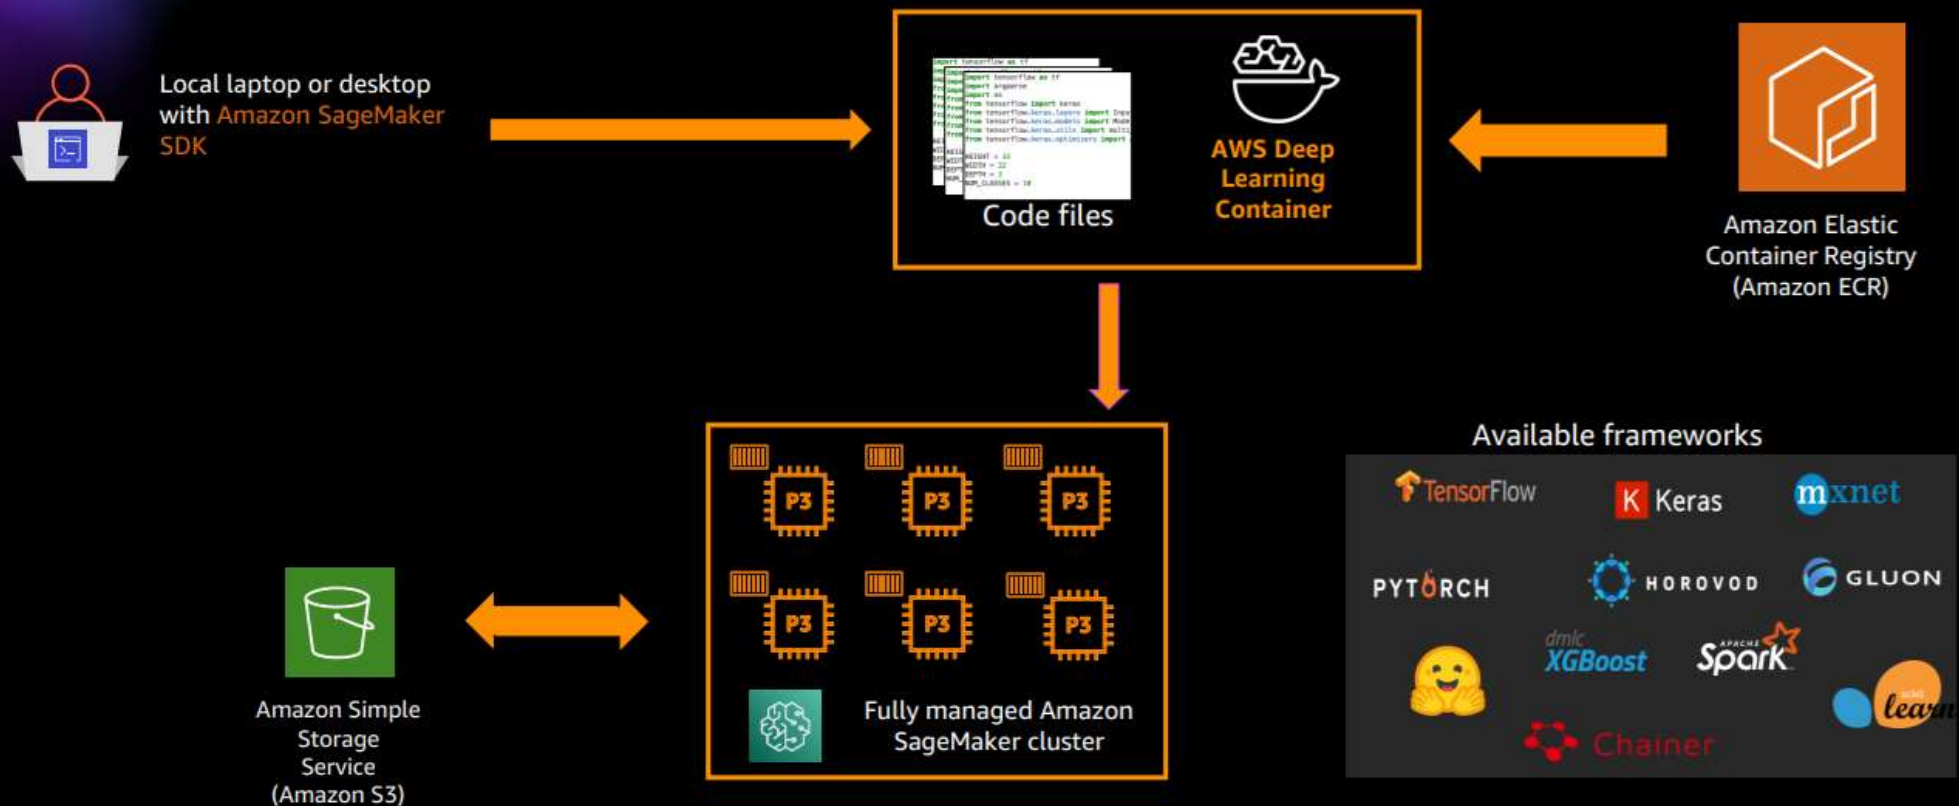
xgb.fit({'train': s3_input_train, 'validation': s3_input_validation})
```



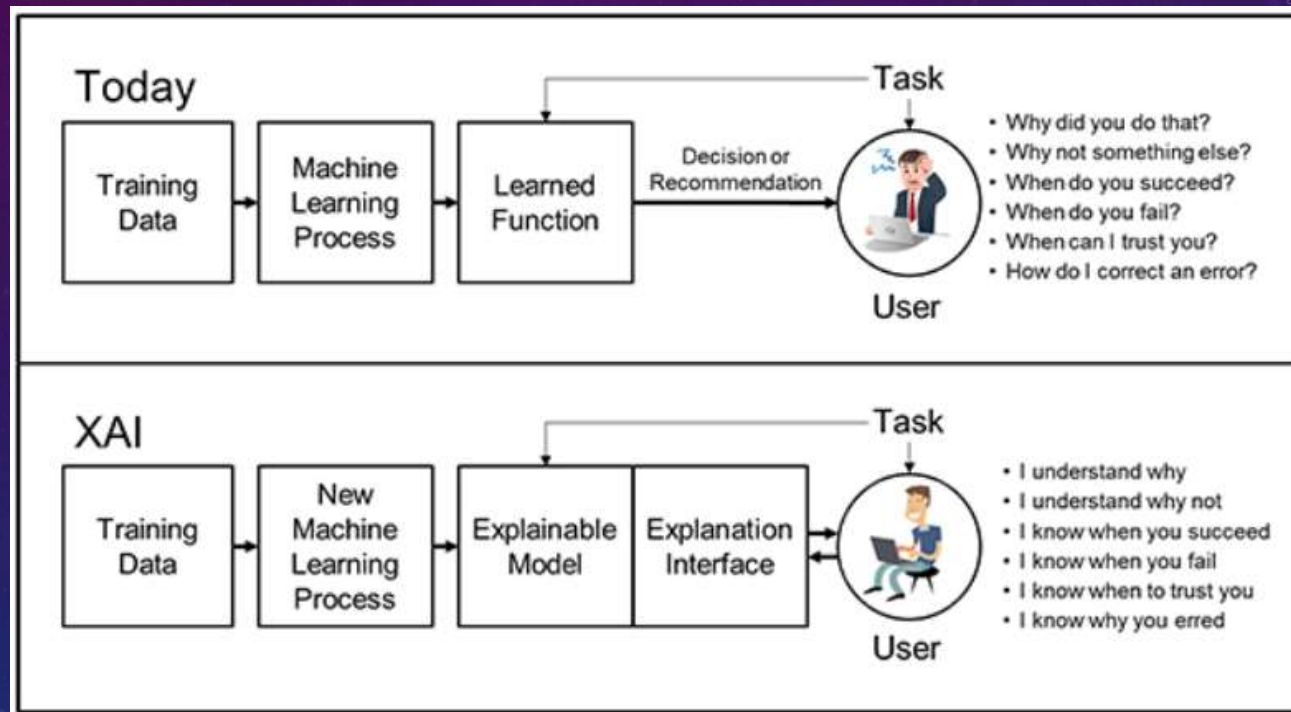
Amazon Elastic Container
Registry
(Amazon ECR)

Bring your own script ('script mode')

HIGH LEVEL WORKFLOW



Understand How Machines Make Decision



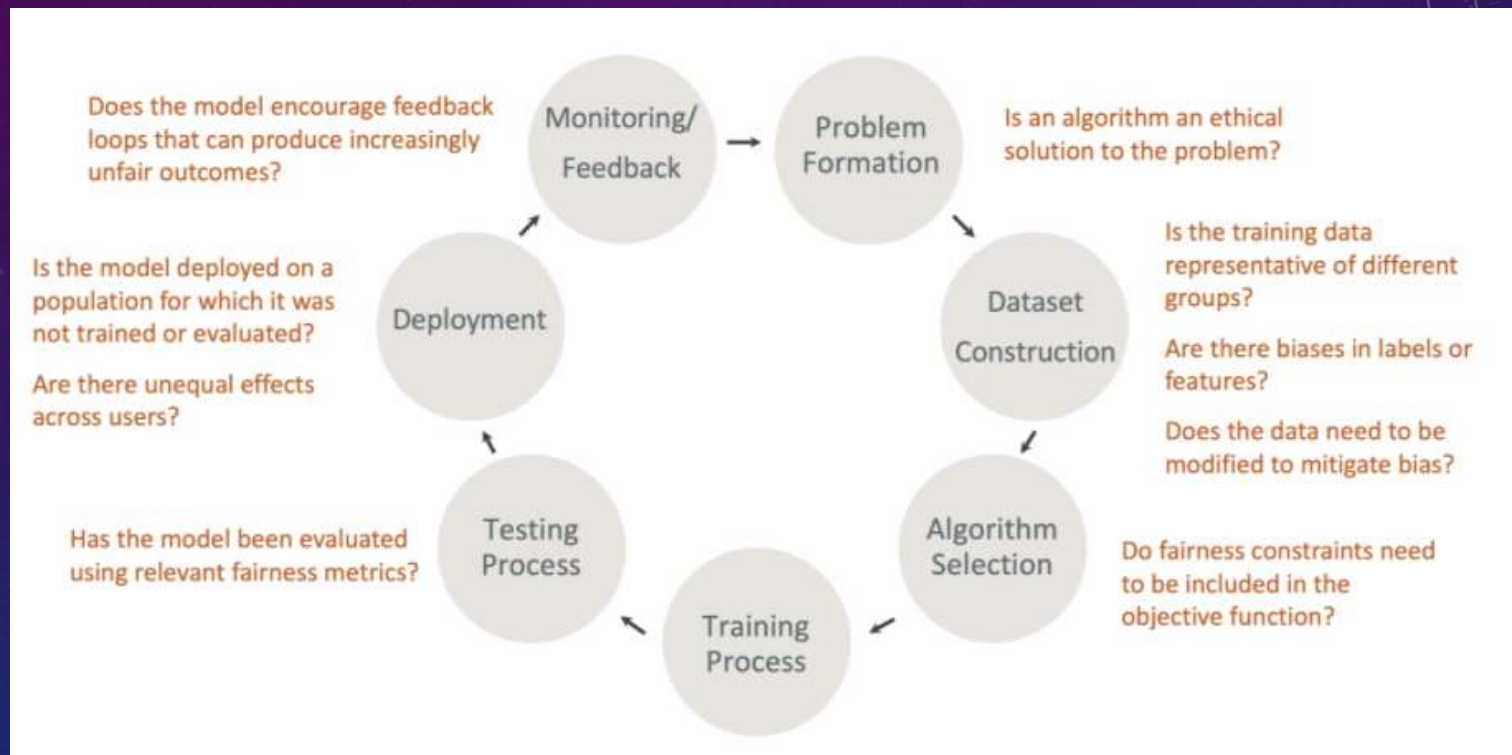
The diagram Source: <https://www.darpa.mil/program/explainable-artificial-intelligence>

Understand How Machines Make Decision

AI-based systems are disrupting almost every industry – healthcare, finances, education, retail, media, marketing, social, economy, HR system, etc, etc. and helping us to make crucial decisions that are impacting millions of lives.

Hence it is important to understand how these decisions are made by the AI system.

Machine Learning Lifecycle



<https://docs.aws.amazon.com/sagemaker/latest/dg/clarify-fairness-and-explainability.html>

Explain the Model Decision

Simple Algorithms

- Bayesian classifiers,
 - Decision trees,
 - Linear Regression,
 - Logistic Regression etc ..
- possess a certain degree of interpretability, traceability, visibility, and transparency in their decision-making process at the cost of the performance

Complex Algorithms

- PCA,
 - Deep Neural Network,
 - Ensemble methods,
 - Random forests etc
- sacrifice their explainability to achieve high performance and accuracy

AWS Sagemaker clarify

- Detecting dataset bias
 - ✓ Pre Training & Post Training Bias Detection
 - ✓ Measure bias using a variety of statistical metrics.
 - ✓ Help in early detection of bias in dataset
- Detecting model bias
 - ✓ Explain how feature values contribute to the predicted outcome, both for the model overall and for individual predictions.
 - ✓ Run a SageMaker Clarify analysis, which includes automatic deployment to a temporary endpoint, and computation of bias metrics using our model and dataset.
 - ✓ By computing these metrics, we can figure out if trained model has similar predictive behavior across groups.

Credit Default
Predicted Rate
= 40%

Guarantor = True

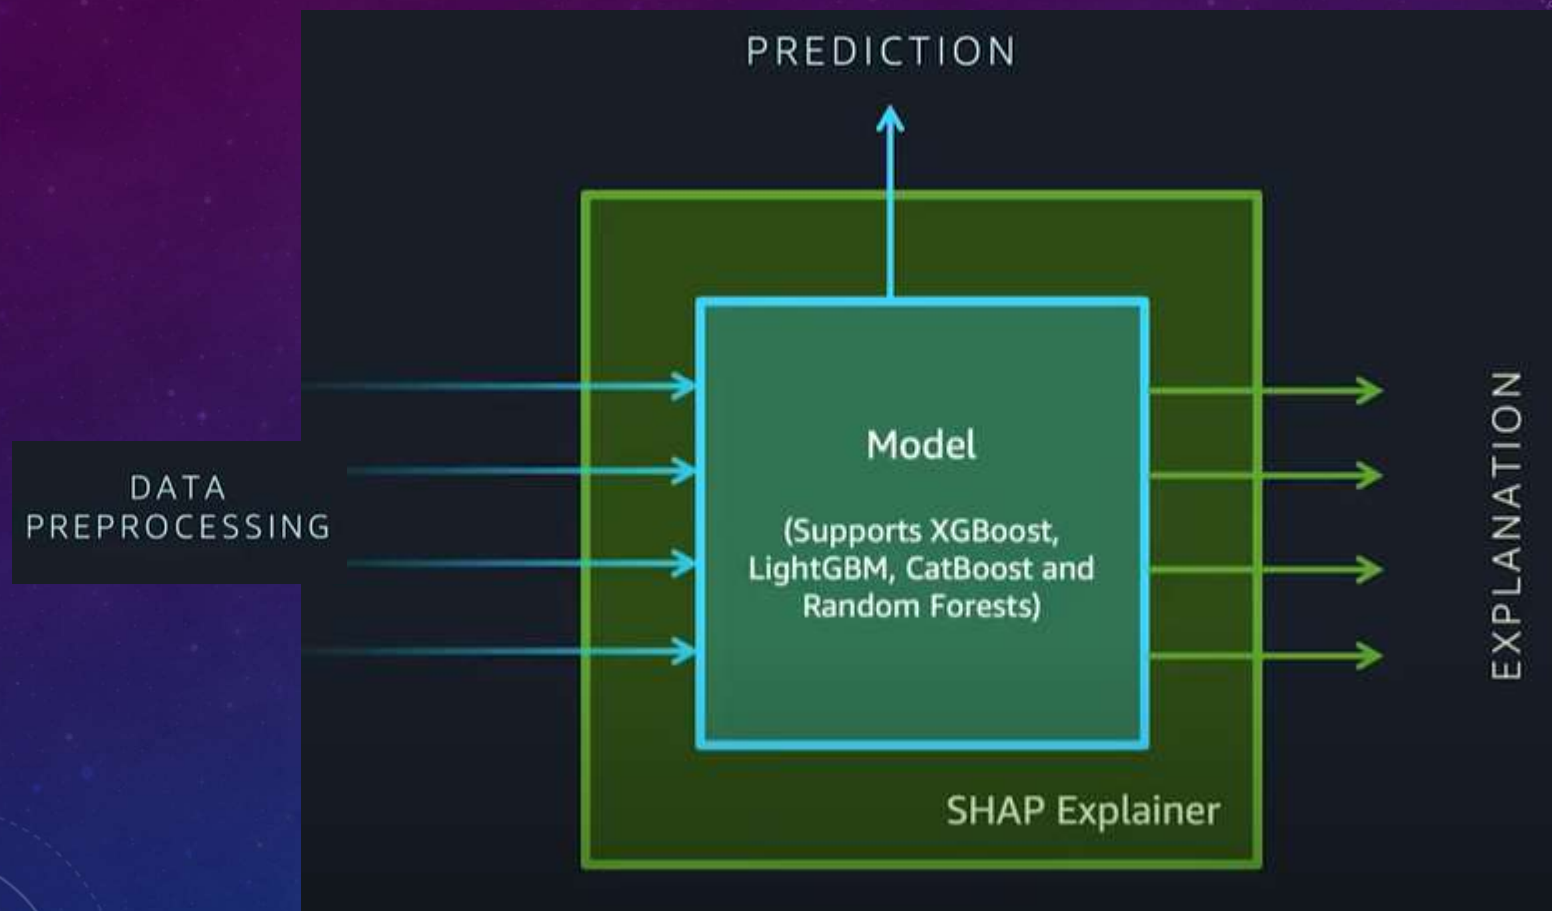
Residence Type = Rental

Age = 18

Credit Amount = \$20,000



Credit Default



SHAP stands for SHapley Additive exPlanations. 'Shapley' relates to a game theoretic concept called Shapley values that is used to create the explanations.

A Shapley value describes the marginal contribution of each 'player' when considering all possible 'coalitions'.

Using this in a machine learning context, a Shapley value describes the marginal contribution of each feature when considering all possible sets of features.

'Additive' relates to the fact that these Shapley values can be summed together to give the final model prediction.

As an example, we might start off with a baseline credit default risk of 10%. Given a set of features, we can calculate the Shapley value for each feature. Summing together all the Shapley values, we might obtain a cumulative value of +30%. Given the same set of features, we therefore expect our model to return a credit default risk of 40% (i.e. $10\% + 30\%$).

Shapley Values

WITH GUARANTOR

Guarantor = True ✓
 Residence Type = Rental ✓
 Age = 18 ✓
 Credit Amount = \$20,000 ✓
 Credit Default Risk = 40%

WITHOUT GUARANTOR

Guarantor = ~~True~~ = False ✗
 Residence Type = Rental ✓
 Age = 18 ✓
 Credit Amount = \$20,000 ✓
 Credit Default Risk = 60%
 Effect of Guarantor = -20%

WITH GUARANTOR

Guarantor = True ✓
 Residence Type = ~~Rental~~ = Own ✗
 Age = 18 ✓
 Credit Amount = \$20,000 ✓
 Credit Default Risk = 10%

WITHOUT GUARANTOR

Guarantor = ~~True~~ = False ✗
 Residence Type = ~~Rental~~ = Own ✗
 Age = 18 ✓
 Credit Amount = \$20,000 ✓
 Credit Default Risk = 20%
 Effect of Guarantor = -10%, -20%

WITH GUARANTOR

Guarantor = True ✓
 Residence Type = ~~Rental~~ = Own ✗
 Age = ~~18~~ = 45 ✗
 Credit Amount = ~~\$20,000~~ = \$5,000 ✗
 Credit Default Risk = 5%

WITHOUT GUARANTOR

Guarantor = ~~True~~ = False ✗
 Residence Type = ~~Rental~~ = Own ✗
 Age = ~~18~~ = 45 ✗
 Credit Amount = ~~\$20,000~~ = \$5,000 ✗
 Credit Default Risk = 45%
 Effect of Guarantor = -10%, -15%, -25%, -5%, -15%, -15%, -10%, -20%

Overall Effect of Guarantor = SHAP Value for Guarantor = -17%

Shapley Values

WITH GUARANTOR

Guarantor = True ✓

Residence Type = ~~Rental~~ = Own ✗

Age = 18 ✓

Credit Amount = ~~\$20,000~~ = \$5,000 ✗

Credit Default Risk = 20%

WITHOUT GUARANTOR

Guarantor = ~~True~~ = False ✗

Residence Type = ~~Rental~~ = Own ✗

Age = 18 ✓

Credit Amount = ~~\$20,000~~ = \$5,000 ✗

Credit Default Risk = 45%

Effect of Guarantor = -25%, -5%, -15%, -15%, -10%, -20%

WITH GUARANTOR

Guarantor = True ✓

Residence Type = ~~Rental~~ = Own ✗

Age = ~~18~~ = 45 ✗

Credit Amount = ~~\$20,000~~ = \$5,000 ✗

Credit Default Risk = 5%

WITHOUT GUARANTOR

Guarantor = ~~True~~ = False ✗

Residence Type = ~~Rental~~ = Own ✗

Age = ~~18~~ = 45 ✗

Credit Amount = ~~\$20,000~~ = \$5,000 ✗

Credit Default Risk = 45%

Effect of Guarantor = -10%, -15%, -25%, -5%, -15%, -15%, -10%, -20%

Overall Effect of Guarantor = SHAP Value for Guarantor = -17%

Explain Credit Decision with Sagemaker Clarify

- Classify credit applications and predict whether the credit would be payed back or not
- Bank to Reduce the risk of losing money due to unpaid credits
- Also reduce the risk of denying trustworthy customers credit which has a set of negative impacts.

AWS Services

As part of the solution, the following services are used:

AWS Lambda: Used to generate a synthetic credits dataset and upload to Amazon S3.

AWS Glue: Used to crawl datasets, and transform the credits dataset using Apache Spark.

Amazon S3: Used to store datasets and the outputs of the AWS Glue Job.

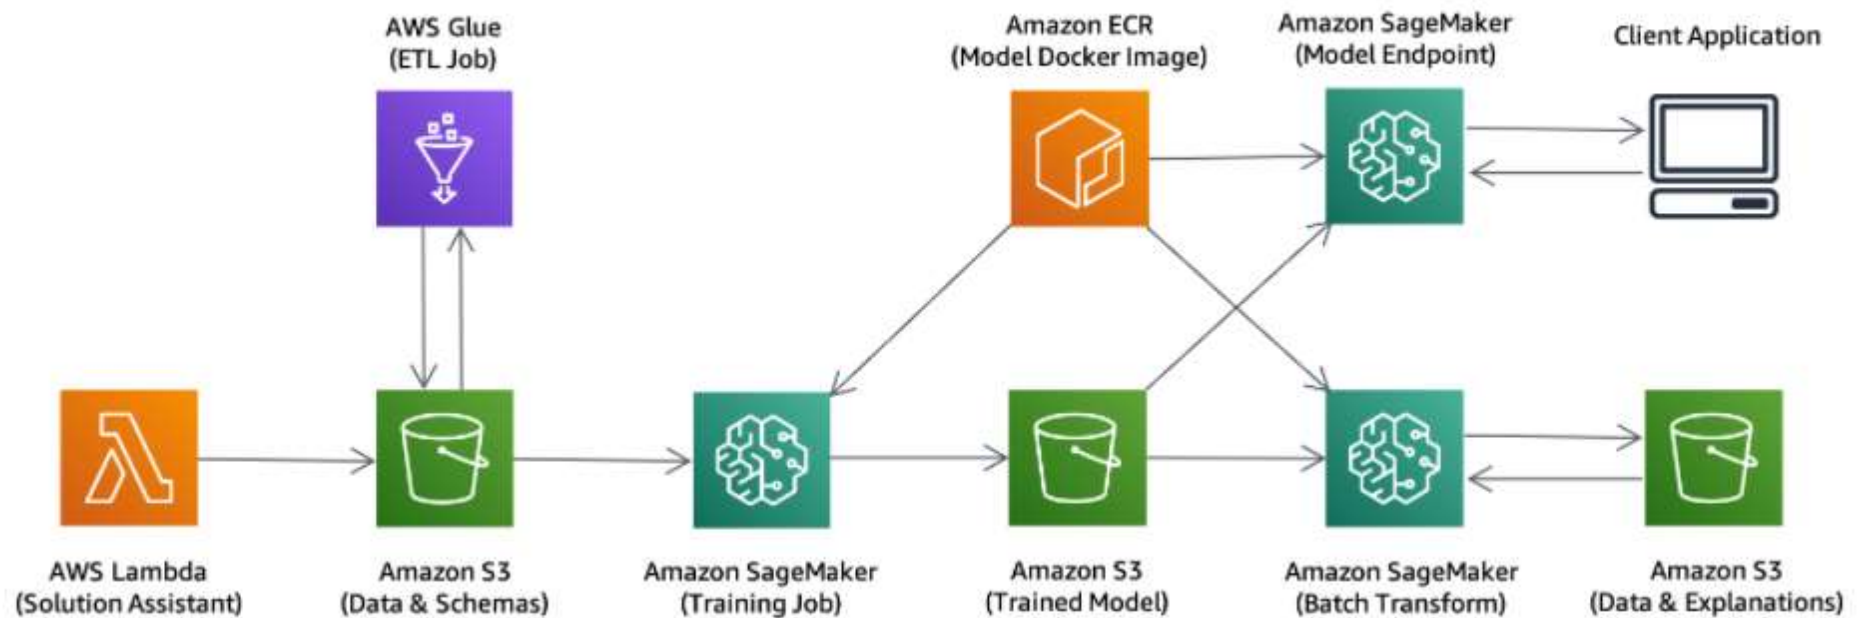
Amazon SageMaker Notebook: Used to train the LightGBM model.

Amazon ECR: Used to store the custom Scikit-learn + LightGBM training environment.

Amazon SageMaker Endpoint: Used to deploy the trained model and SHAP explainer.

Amazon SageMaker Batch Transform: Used to compute explanations in batch.

Architecture



AWS Sagemaker Studio

The screenshot displays the AWS Sagemaker Studio interface. On the left, a sidebar titled 'SageMaker resources' contains a 'Projects' dropdown menu, a search bar, and a 'Create project' button. Below this, a list of resources is shown, with 'demosagemakerstudio' selected. The main panel features a tabbed interface with 'SageMaker JumpStart' and 'Explain Credit Decisions' tabs. The 'Explain Credit Decisions' tab is active, showing the solution's name, deployment status ('Deployed 11 days ago'), and a 'Browse JumpStart' button. The main content area is divided into two sections: 'Open solution in Studio' with an 'Open Notebook' button, and 'Delete solution' with a 'Delete all resources' button.

SageMaker resources
Select the resource to view.

Projects

Search

Create project

Name

demosagemakerstudio

End of the list

SageMaker JumpStart × Explain Credit Decisions ×

Explain Credit Decisions

Browse JumpStart

Deployed 11 days ago

Base Solution: [Explain Credit Decisions](#)

Open solution in Studio

Open a notebook to walk through the solution using the data provided or customize it to use your own.

Open Notebook

Delete solution

If you no longer want this solution you can easily delete the resources it created. Selecting the button below will automatically delete all standard resources that were created when launching the solution, including the S3 bucket.

Delete all resources

Stages

- Our solution is split into the following stages, and each stage has its own notebook:
- Introduction: We take a high-level look at the solution components.
- Datasets: We prepare a dataset for machine learning using AWS Glue.
- Training: We train a LightGBM model using Amazon SageMaker, so we have an example trained model to explain.
- Endpoint: We deploy the model explainer to a HTTP endpoint using Amazon SageMaker and visualize the explanations.
- Batch Transform: We use Amazon SageMaker Batch Transform to obtain explanations for our complete dataset.
- Dashboard: We develop a dashboard for explanations using Amazon SageMaker and Streamlit.
- Conclusion: We wrap things up and discuss how to clean up the solution.

Model Deployment

- Once the SageMaker training job is completed, a number of trained model artifacts are stored in S3 bucket
- Retrieve the model data (i.e. model.tar.gz) from the most recent trained model
- Define the model using SKLearnModel() to deploy which includes the explainer logic.
- Calling model.deploy() will start a container to host the model.
- Entities are used ->

```
entities = [  
    'data',  
    'features',  
    'descriptions',  
    'prediction',  
    'explanation_shap_values',  
    'explanation_shap_interaction_values'  
]
```

Model Explanation


- Call `explainer.predict` with features (for a credit application) to obtain a prediction and explanation.
- Visualize Explanations with bokeh & waterfall chart.
- A waterfall chart can be used to show the cumulative effect of each feature.
- Green arrows indicate that the feature *decreased* the predicted credit default risk for the individual credit application.
- While red arrows indicate that the feature *increased* the predicted credit default risk for the individual credit application.
- After all features have been considered, we reach the final predicted credit default risk (at the top of the chart).



- Home
- Data
- AutoML
- Experiments
- Notebook jobs
- Pipelines
- Models
- Deployments
- Quick start solutions
 - Solutions, models, example notebooks
 - Pretrained models, example notebooks, & prebuilt solutions
 - Launched Quick start assets
 - Manage launched Quick start solutions
- Learning resources
 - Documentation

we then show the associated waterfall chart.

```
[16]: x_axis_label = 'Credit Default Risk Score (%)'
summary_waterfall = visuals.WaterfallChart(
    baseline=explanation_summary['expected_value'],
    shap_values=explanation_summary['shap_values'],
    names=explanation_summary['feature_names'],
    descriptions=explanation_summary['feature_descriptions'],
    max_features=10,
    x_axis_label=x_axis_label,
)
summary_waterfall.show()
```




We can see from the summary waterfall chart above that features related to finance have the largest combined effect on the credit default risk. Although features related to finance reduce the credit default risk, the features related to employment bring the risk back up again to a certain degree.

Model Explanation

```
x_axis_label = 'Credit Default Risk Score (%)'
summary_waterfall = visuals.WaterfallChart(
    baseline=explanation_summary['expected_value'],
    shap_values=explanation_summary['shap_values'],
    names=explanation_summary['feature_names'],
    descriptions=explanation_summary['feature_descriptions'],
    max_features=10,
    x_axis_label=x_axis_label,
)
summary_waterfall.show()
```

We can see from the summary waterfall chart that features related to finance have the largest combined effect on the credit default risk.

Although features related to finance reduce the credit default risk, the features related to employment bring the risk back up again to a certain degree.



Explain Prediction

```

detailed_waterfall = visuals.WaterfallChart(
    baseline=explanation['expected_value'],
    shap_values=explanation['shap_values'],
    names=explanation['feature_names'],
    feature_values=explanation['feature_values'],
    descriptions=explanation['feature_descriptions'],
    max_features=10,
    x_axis_label=x_axis_label
)
detailed_waterfall.show()

```



Explain Prediction

- Detailed waterfall chart shows that
 - Not having a checking account with the same bank indicates a lower credit default risk.
 - The credit to purchase a used car is associated with a lower credit default risk
 - After this we see a number of features that increase the credit default risk:
 - A credit amount of 6000 EUR, a lack of employment and a credit duration of 36 months.
 - Another potential area for investigation, would be related to the repayment history feature.
 - Not having a very poor repayment history is associated with a higher credit default risk score.

Counterfactual Example

```
counter_sample = dict(sample)
counter_sample['finance__accounts__checking__balance'] = 'negative' # from 'no_account'
counter_output = explainer.predict(counter_sample)
counter_explanation = visuals.detailed_explanation(counter_output)
visuals.WaterfallChart(
    baseline=counter_explanation['expected_value'],
    shap_values=counter_explanation['shap_values'],
    names=counter_explanation['feature_names'],
    feature_values=counter_explanation['feature_values'],
    descriptions=counter_explanation['feature_descriptions'],
    max_features=10,
    x_axis_label=x_axis_label,
).show()
```



Counterfactual Example

- Now let's switch the value of the checking account balance of the applicant from no account to negative.
- We can then see how the overall prediction of the model changes, and also see the updated contribution of this feature.
- Clearly, this application has become substantially more risky.

Home

Shared models

Manage shared models & notebooks

Inference compiler

Optimize trained models

Edge packager

Package models for edge devices

Deployments

Projects

Automate model building & deployment

Inference recommender

Optimize inference cost & performance

Endpoints

Manage deployed models

Fleets & devices

Manage edge device groups

Quick start solutions

Solutions, models, example notebooks

Pretrained models, example notebooks, & prebuilt solutions

Laun X Home X 0_int X 1_dat X 2_tra X 3_en X Pipel X 4_bal X Expe X Proje X Endp X sager X Desci X 5_cor X

Endpoints

Show introduction

Search...

Name	ARN	Created	Status	Last modified
sagemaker-soln-ecl-j-e0e4av-...	arn:aws:sagemaker:us-east-1:396548483...	34 minutes ago	In service	32 minutes ago
End of results				

Amazon SageMaker Studio

File Edit View Run Kernel Git Tabs Settings Help

default1-16nov2022 / Perso

Manage shared models & notebooks

Inference compiler
Optimize trained models

Edge packager
Package models for edge devices

Deployments

Projects
Automate model building & deployment

Inference recommender
Optimize inference cost & performance

Endpoints
Manage deployed models

Fleets & devices
Manage edge device groups

Quick start solutions

Solutions, models, example notebooks
Pretrained models, example notebooks, & prebuilt solutions

less than 10 seconds ago

ENDPOINT DETAILS

sagemaker-soln-ecd-js-e0e4av-explainer

Test inference Data quality Model quality Model explainability Model bias Monitoring job history AWS settings

Test inference

Test and validate your model by sending a request to the SageMaker hosted endpoint and receiving a response from your model.

JSON editor

1 {

2

3 }

> Configure endpoint URL and headers -
optional

Send Request

DEMO

References

<https://github.com/aws-labs/sagemaker-explaining-credit-decisions>

https://sagemaker-examples.readthedocs.io/en/latest/sagemaker_processing/fairness_and_explainability/fairness_and_explainability.html

My Github link : <https://bit.ly/3BzAsQT>

Linkedin : <https://www.linkedin.com/in/sarbani-maiti-35b89111/>

THANK YOU

